



# The Well-Dressed Bioinformatician

(NGS  
Computational  
Methods)

I got some sequence data. What now?



- I. File Formats**
- II. Quality Control
- III. Sequence Alignment
- IV. Motif Discovery
- V. De Novo Genome  
Assembly
- VI. ChIP-seq
- VII. RNA-seq

# Some common file formats

- FASTA
  - Just the sequence, ma'am
- FASTQ
  - FASTA + “Quality” score for each base.
  - Quality is ASCII (33 + Phred score)

```
@SEQ_ID
```

```
GATTTGGGGTTCAAAGCAGTATCGATCAAATAGTAAATCCATTTGTTCAACTCACAGTTT  
+  
!''*(((((***+))%%%++)(%%%%).1***-+*''))**55CCF>>>>>CCCCCCC65
```




Figure 1. An example of a DNA sequence tracing and the Phred score (grey bars) corresponding to each colored peak. The colored peaks on the trace correspond to each DNA letter. For example 'T' bases are represented in red, and this sequence has four 'T' bases on a row, as viewed by the four red peaks in the sequence. The aqua horizontal line placed across the grey bars represents a Phred score of 20 which is considered an acceptable level of accuracy. As indicated in Table 1, a Phred score of 20 corresponds to a 99% accuracy in the base call. Therefore, bars above this line indicate base calls that have a higher than 99% probability of being correct. Those below have less than a 99% probability of being correct. Sequence tracing program is courtesy of FinchTV ([www.geospiza.com](http://www.geospiza.com)).

# Phred scores

Q = Phred score

P = error probability

$$Q = -10 \log_{10} P$$

$$P = 10^{\frac{-Q}{10}}$$

# More file formats

- SRA
  - native format for Sequence Read Archive (repository of NGS data)
  - Aligned sequences can be compressed by genome reference
  - Convert to FASTQ with **sra-toolkit** (fastq-dump)

# More file formats

- SAM
  - **S**equence **A**lignment / **M**ap
  - generic format for aligned sequences
- BAM
  - binary version of SAM
  - smaller
  - most software wants this
- **Samtools**
  - Convert BAM to SAM
  - Sort and index BAMs
  - Count number of reads at a position (pileup)
  - pysam: a python wrapper
  -

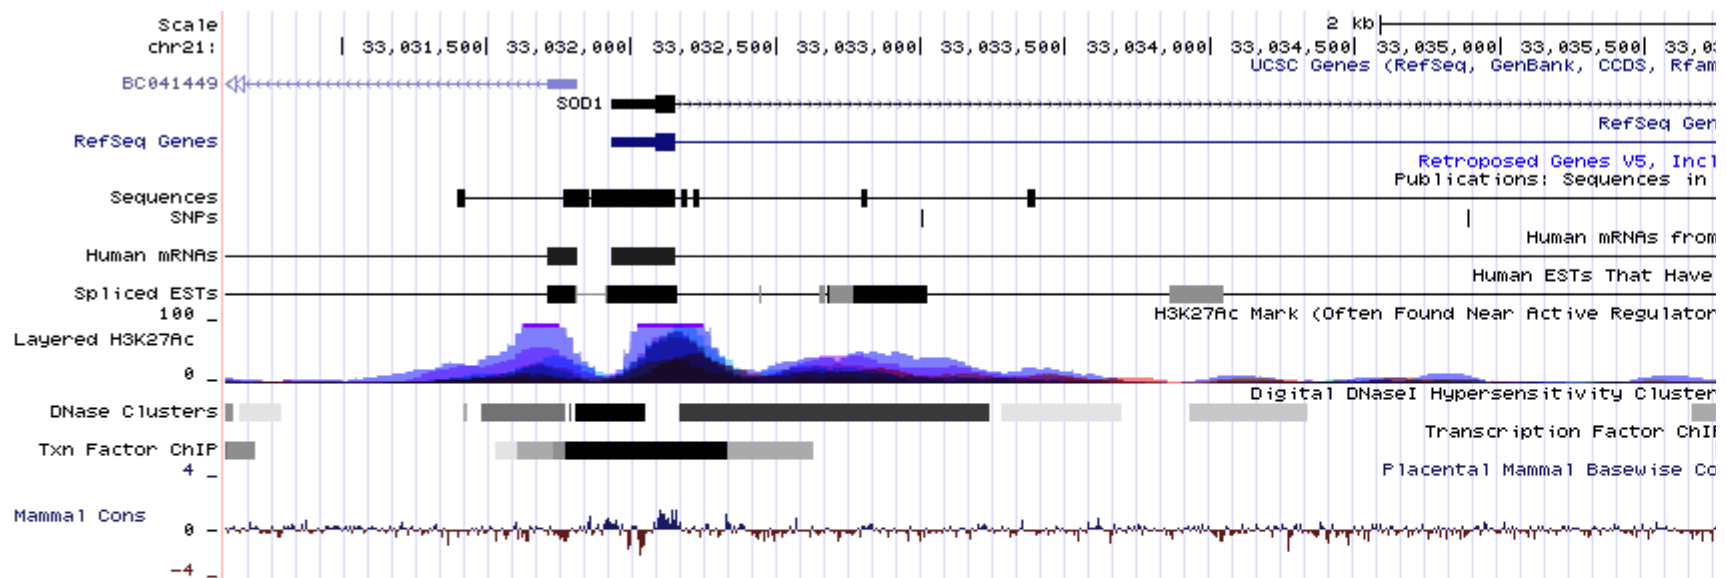


# More file formats

- BED
  - Discrete genome annotations (regions)
    - Gene positions
    - Exons / Introns
    - Transcription start sites
    - Protein binding sites
    - Histone locations
    - etc.
  - Contains:
    - Chromosome
    - Start position
    - Stop position
    - Description fields
- **Bedtools**
  - GFF and GTF are similar formats.
  - All are tab-separated, human-readable

# More file formats

- wig and bigWig
  - Wiggle tracks
  - Continuous-valued genome annotations






- I. File Formats
- II. Quality Control**
- III. Sequence Alignment
- IV. Motif Discovery
- V. De Novo Genome  
Assembly
- VI. ChIP-seq
- VII. RNA-seq


# FASTQC

- Graphical tool to evaluate NGS quality
- Can be run in interactive or command line mode
- Generates HTML reports
- Evaluates results:




**Note: Depending on your experiment, it may give false alarms**


# Per base sequence quality: good




# Per base sequence quality: bad



# Per sequence quality scores: good




# Per sequence quality scores: bad







# Per base sequence content: good




# Per base sequence content: bad




# Per base GC content: good




# Per base GC content: bad



# Per sequence GC content: good



# Per sequence GC content: bad





- I. File Formats
- II. Quality Control
- III. Sequence Alignment**
- IV. Motif Discovery
- V. De Novo Genome  
Assembly
- VI. ChIP-seq
- VII. RNA-seq

# Sequence Alignment

- Alignment with exact match

ABCTUV  
ABUV

ABCTUV  
AB--UV



# Sequence Alignment

- Inexact

ATTCCATTCATAA  
ATTCTTAATAA

ATTCCATTCATAA  
ATTC--TTAATAA

# Sequence Alignment

- Global alignment

CATGCTATGGTCA  
ATGCGGTCGACAC

- Local alignment

CATGCTATGGTCA  
ATGC---GGTC

# Pairwise Alignment Algorithms

- Exact (returns optimal solution)
  - Needleman – Wunsch (1970)
    - Global
  - Smith-Waterman (1981)
    - Local
- Heuristic (trade exactness for speed)
  - FASTA (1985)
  - BLAST (1990)
  - MUMmer (1999)
  - Vmatch
  - Many more!

# Multiple Alignment Algorithms

- Optimal solution:
  - $O(\text{length}^{\text{num\_sequences}})$
- Approximations:
  - ClustalW, T-coffee
    - Progressively combining pairwise alignments
  - SAM, HMMER
    - Hidden Markov models
  - SAGA, RAGA
    - Genetic algorithms
  - MSASA
    - Simulated annealing

# Short Read Alignment

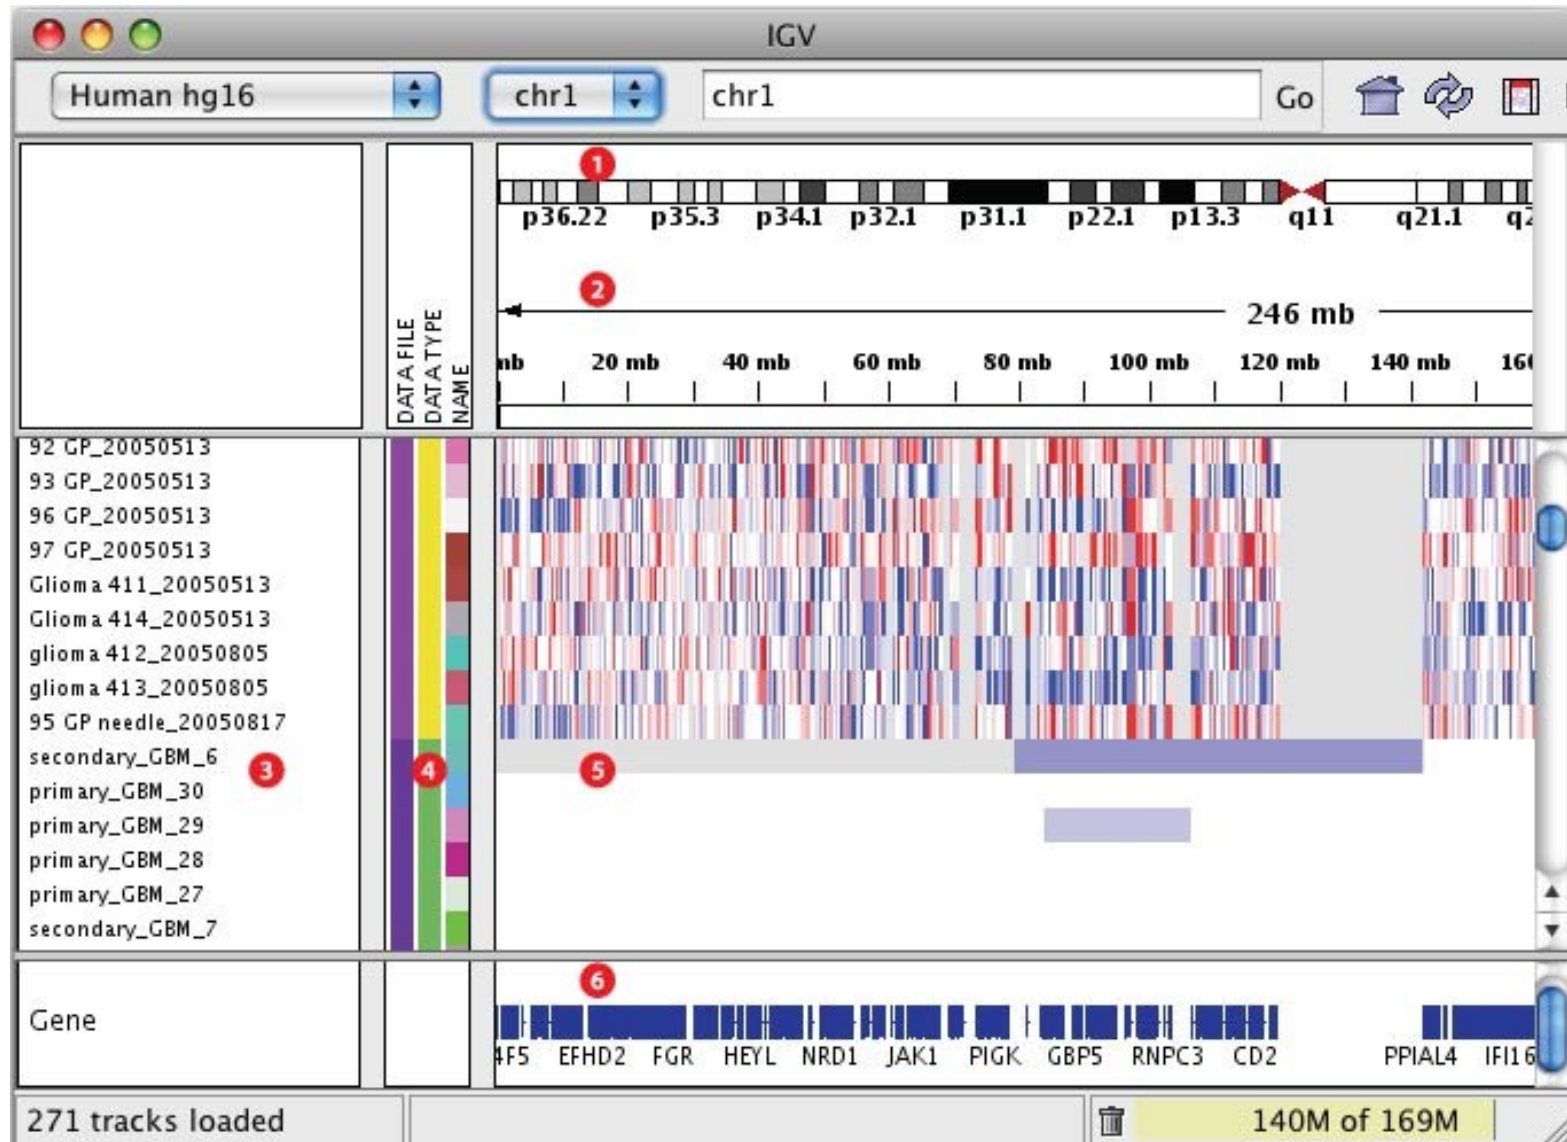
- Some ways to improve performance:
  - Compress and index the genome
    - k-mer hashing
    - Suffix trees
    - Enhanced suffix arrays
    - Burrows-Wheeler transform (BWT)
  - Seed local alignments with exact matches

# Short Read Alignment

- BFAST
  - Explicit time vs. accuracy tradeoff
  - Handles indels
- BWA
  - Handles indels
- Bowtie
  - Doesn't handle indels
- Bowtie2
  - Handles indels
  - Supports local alignment
  - Handles paired end sequencing better



# IGV





- I. File Formats
- II. Quality Control
- III. Sequence Alignment
- IV. Motif Discovery**
- V. De Novo Genome  
Assembly
- VI. ChIP-seq
- VII. RNA-seq



# Motif discovery (*de novo*)



# Motif discovery (*de novo*)




- HOMER
  - **H**ypergeometric **O**ptimization of **M**otif **E**n**R**ichment
  - Motif analysis & more!
  - Chris Brenner (Glass lab)

A

	1	2	3	4	5	6	7
A	1	4	1	2	0	17	13
C	28	5	5	0	3	3	2
G	0	0	4	0	25	1	7
T	2	22	21	29	4	10	9

B





- I. File Formats
- II. Quality Control
- III. Sequence Alignment
- IV. Motif Discovery
- V. De Novo Genome  
Assembly**
- VI. ChIP-seq
- VII. RNA-seq

# *De novo* genome assembly

- Traditionally done with Sanger sequencing
  - Long reads ~1000s of bp
  - Tools (overlap-layout-consensus)
    - Atlas
    - ARACHNE
    - Celera
    - PCAP
    - *phrap*
    - Phusion

# *De novo* genome assembly

- Overlap-layout-consensus approach
  - Does OK with small numbers of long reads
  - poorly suited to millions of short NGS reads
    - 10s to 100s of bp
  - Doesn't handle repeats well – so generally just ignores them



# An Eulerian path approach to DNA fragment assembly

Pavel A. Pevzner\*, Haixu Tang<sup>†</sup>, and Michael S. Waterman<sup>†‡§</sup>

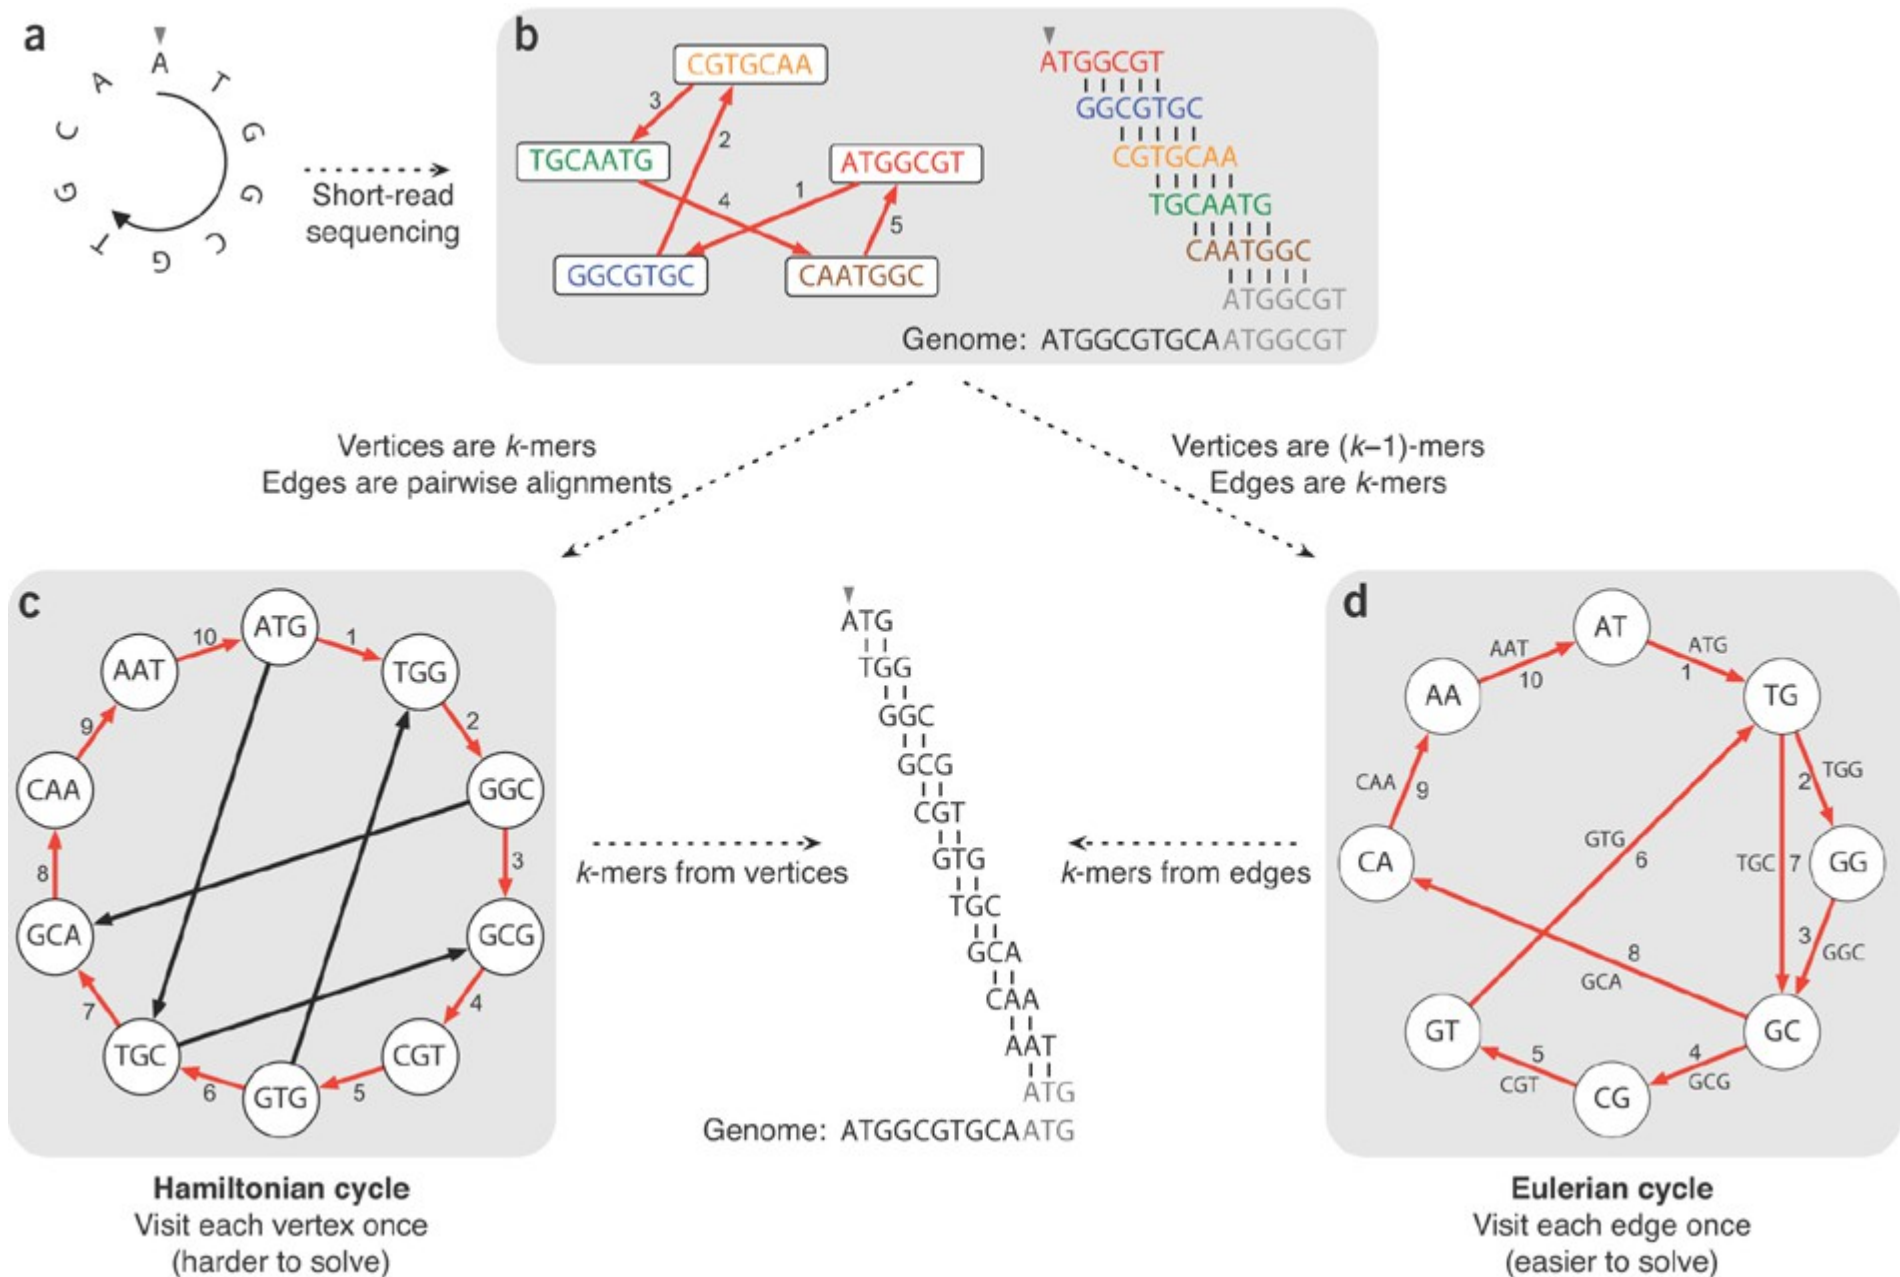
\*Department of Computer Science and Engineering, University of California, San Diego, La Jolla, CA; and Departments of <sup>†</sup>Mathematics and

<sup>‡</sup>Biological Sciences, University of Southern California, Los Angeles, CA

Contributed by Michael S. Waterman, June 7, 2001

For the last 20 years, fragment assembly in DNA sequencing followed the “overlap–layout–consensus” paradigm that is used in all currently available assembly tools. Although this approach proved useful in assembling clones, it faces difficulties in genomic shotgun assembly. We abandon the classical “overlap–layout–consensus” approach in favor of a new EULER algorithm that, for the first time, resolves the 20-year-old “repeat problem” in fragment assembly. Our main result is the reduction of the fragment assembly to a variation of the classical Eulerian path problem that allows one to generate accurate solutions of large-scale sequencing problems. EULER, in contrast to the CELERA assembler, does not mask such repeats but uses them instead as a powerful fragment assembly tool.

Because the Eulerian path approach transforms a once difficult layout problem into a simple one, a natural question is: “Could the Eulerian path approach be applied to fragment assembly?” Idury and Waterman, mimicked fragment assembly as an SBH problem (11) by representing every read of length  $n$  as a collection of  $n - l + 1$  overlapping  $l$ -tuples (continuous short strings of fixed length  $l$ ). At first glance, this transformation of every read into a collection of  $l$ -tuples (breaking the puzzle into smaller pieces) is a very short-sighted procedure, because information about the sequencing reads is lost. However, the loss of information is minimal for large  $l$  and is well paid for by the computational advantages of the Eulerian path approach. In addition, lost information can be restored at later stages.

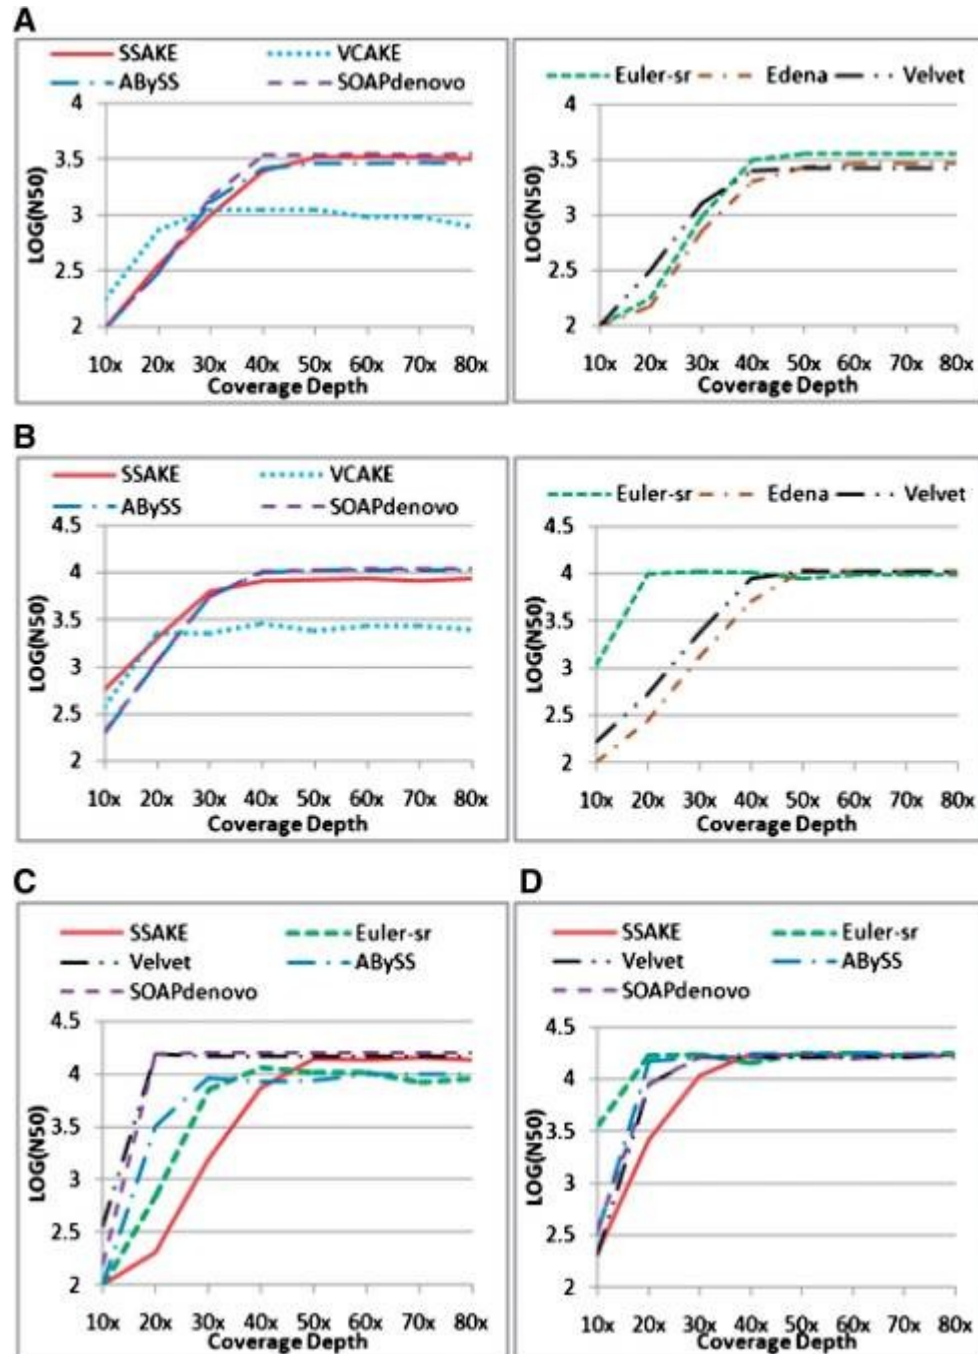




# Velvet

- `velveth`
  - Creates hashtable for `velvetg`
  - May want to play with k-mer length (hash length)
- `velvetg`
  - Assembles using de Bruijn graphs
- `VelvetOptimizer`
  - optimizes parameters



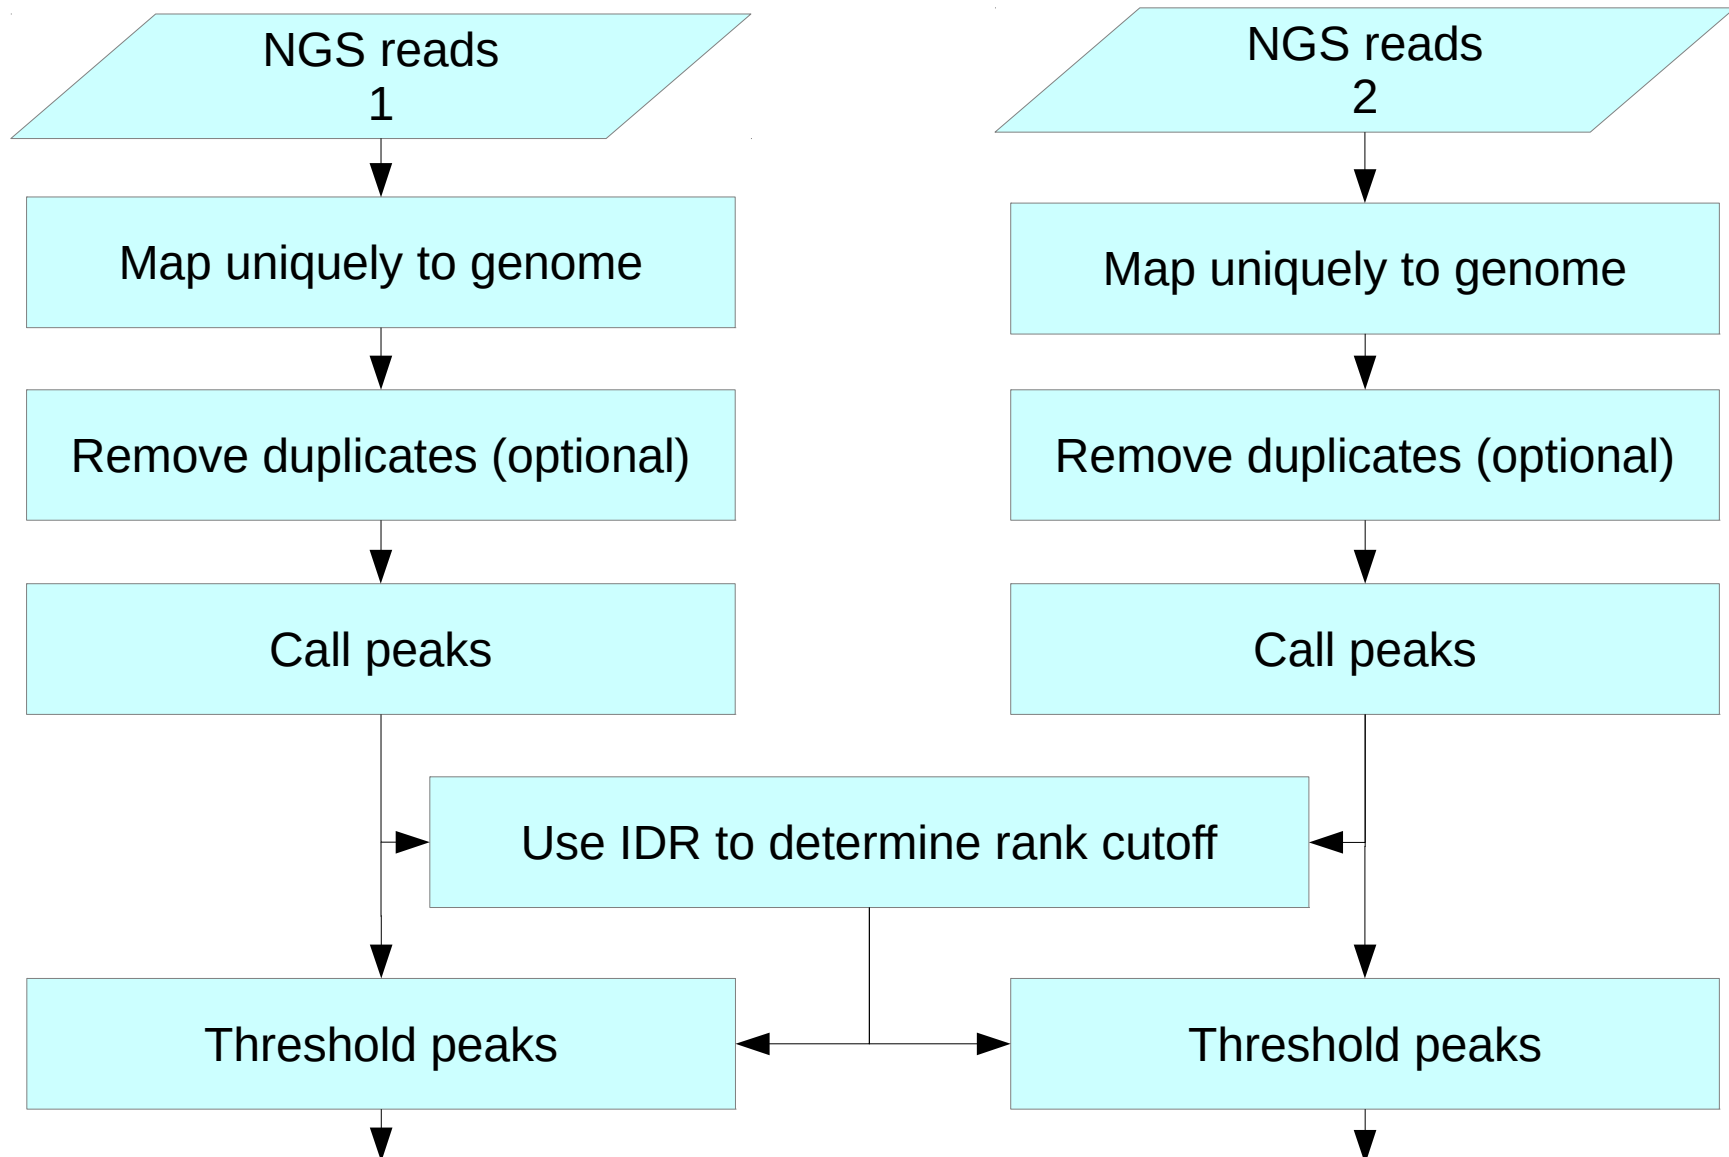


Comparison of the effect of various coverage depths on N50 length in T.bru assembly when BCER was 0.6%. (A) Single-end reads assembly, read length (RL)=35 bp; (B) single-end assembly, RL=75 bp; (C) paired-end reads assembly, RL=35 bp; (D) paired-end assembly, RL=75 bp.



- I. File Formats
- II. Quality Control
- III. Sequence Alignment
- IV. Motif Discovery
- V. De Novo Genome  
Assembly
- VI. ChIP-seq**
- VII. RNA-seq


# ChIP-seq data analysis



# Duplicate removal

- PCR amplification step produces multiple copies of same fragment
  - sequence composition, size, etc. biases distribution
  - shows up as multiple reads with same 5' position
- However, expect some proportion of “real” duplicate reads based on coverage
- Benefits of removal for various applications debated but most favor removal

# Peak Calling




**nrsf / Peaks detected by PICS (Region List)**

143,411

**Genes / Gene Annotations (Genes)**

KCNS2

chr 8



# Peak Calling

- Large (40+) and expanding panoply of algorithms
- “It is challenging to rigorously assess subtle improvements due to the scarcity and unreliability of verified binding sites for any ChIP-seq dataset.”
  - Wilbanks and Facciotti (2010)



# ChIP-seq peak calling programs selected for evaluation in Wilbanks and Facciotti (2010)

Program	Reference	Version	Graphical user interface?	Window-based scan	Tag clustering	Gaussian kernel density estimator	Strand-specific density	Peak height or fold enrichment (FE)	Background subtraction	Compensates for genomic duplications or deletions	False Discovery Rate	Compare to normalized control data (FE)	Compare to statistical model fitted with control data	Statistical model or test
CisGenome	28	1.1	X*	X			X	X		X		X		conditional binomial model
Minimal ChipSeq Peak Finder	16	2.0.1			X		X				X			
E-RANGE	27	3.1			X		X				X	X		chromosome scale Poisson dist.
MACS	13	1.3.5		X			X			X		X		local Poisson dist.
QuEST	14	2.3			X		X			X**		X		chromosome scale Poisson dist.
HPeak	29	1.1		X			X					X		Hidden Markov Model
Sole-Search	23	1	X	X			X		X			X		One sample t-test
PeakSeq	21	1.01			X		X					X		conditional binomial model
SISSRS	32	1.4		X			X				X			
spp package (wtd & mtc)	31	1.7		X			X		X	X'	X			
				Generating density profiles			Peak assignment		Adjustments w. control data		Significance relative to control data			

X\* = Windows-only GUI or cross-platform command line interface

X\*\* = optional if sufficient data is available to split control data

X' = method excludes putative duplicated regions, no treatment of deletions

Wilbanks EG, Facciotti MT (2010) Evaluation of Algorithm Performance in ChIP-Seq Peak Detection. PLoS ONE 5(7): e11471.

doi:10.1371/journal.pone.0011471

<http://www.plosone.org/article/info:doi/10.1371/journal.pone.0011471>



# Rapid innovation in ChIP-seq peak-calling algorithms is outdistancing benchmarking efforts

Adam M. Szalkowski and Christoph D. Schmid

Submitted: 2nd August 2010; Received (in revised form): 8th October 2010

## Abstract

The current understanding in the determination of chromatin interactions (ChIP-seq) provides a foundation for peak-calling software. Such loci of peak-calling software peak-callers have been used to generate results. Yet, peak-callers and explain potential

**A manually curated ChIP-seq benchmark demonstrates room for improvement in current peak-finder programs**

**Morten Beck Rye<sup>1,\*</sup>, Pål Sætrom<sup>1,2</sup> and Finn Drabløs<sup>1</sup>**

**Keywords:** ChIP-seq

<sup>1</sup>Department of Cancer Research and Molecular Medicine, Norwegian University of Science and Technology, NO-7489 Trondheim and <sup>2</sup>Department of Computer and Information Science, Norwegian University of Science and Technology, NO-7491 Trondheim, Norway

Received June 24, 2010; Revised October 21, 2010; Accepted November 3, 2010

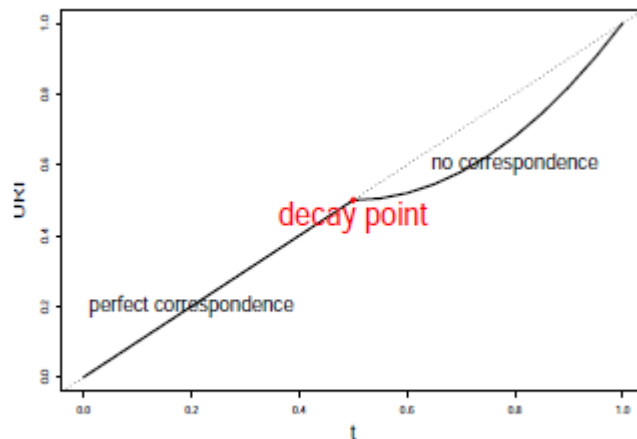
The MACS methodology initially identifies a set of high-confidence peaks which exhibit between 20-100 fold enrichment in reads and then uses these to model the average distance between reads on the two DNA strands.

It then uses a binomial distribution to remove duplicate tags which are represented at statistically unlikely levels within the sample.


Next it generates a full set of candidate peaks by comparing the enrichment within windows to a Poisson distribution generated from the control (input DNA or mock IP)

# Irreproducible discovery rate (IDR)

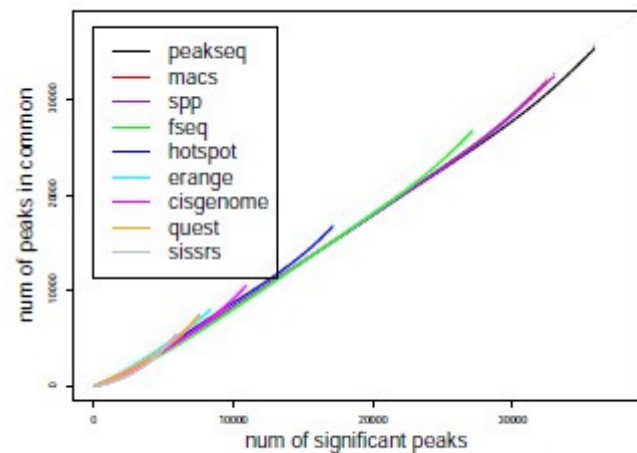
Li, Q., Brown, J.B., Huang, H., and Bickel, P.J. (2011). Measuring reproducibility of high-throughput experiments. The Annals of Applied Statistics 5, 1752–1779.



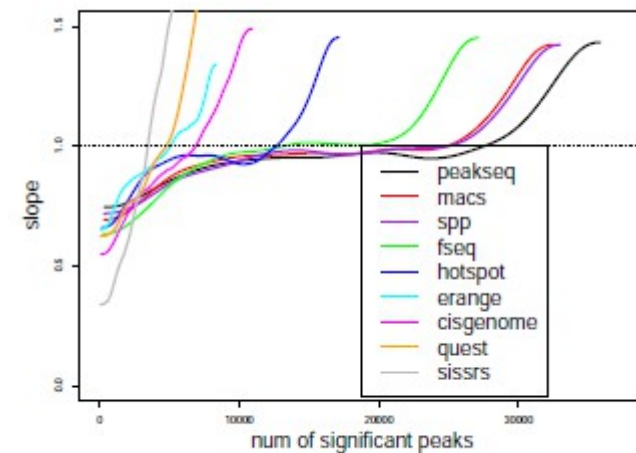
a.  $\Psi$



b.  $\Psi'$

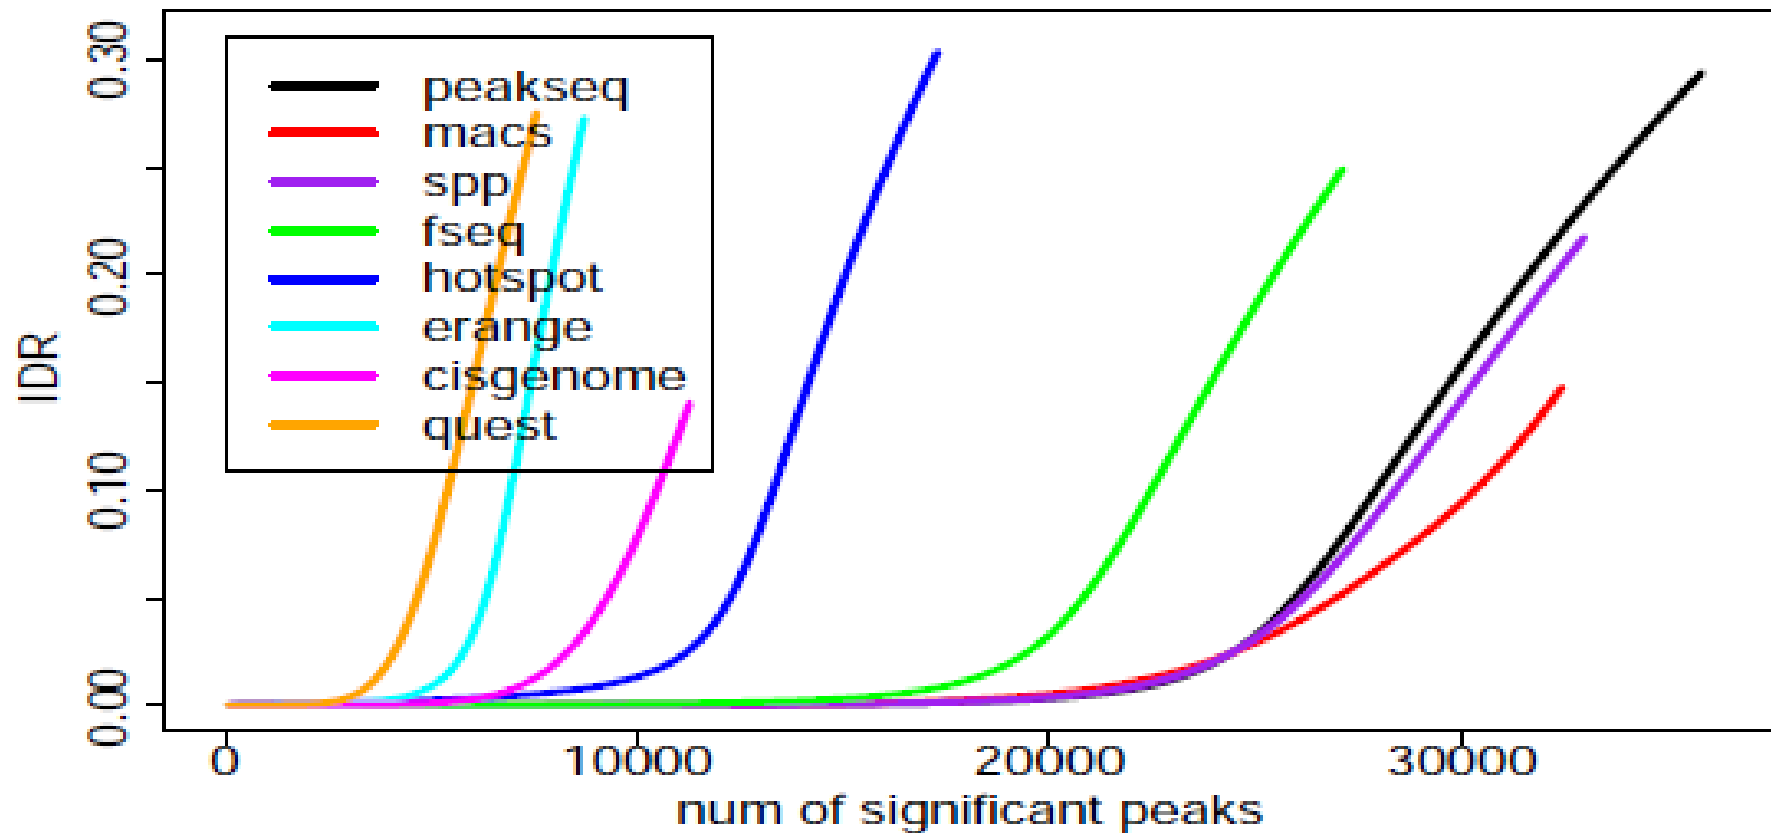


$\Psi$



$\Psi'$

# Irreproducible discovery rate (IDR)





- I. File Formats
- II. Quality Control
- III. Sequence Alignment
- IV. Motif Discovery
- V. De Novo Genome  
Assembly
- VI. ChIP-seq
- VII. RNA-seq**

# RNA-seq

- Cufflinks
  - Analyze transcriptomes (at least 2 conditions) for differential expression & significance



# RNA-seq

- CummeRbund
  - R package
  - Visualizing results from cuffdiff
  - Create SQLite database of results & relationships



# RNA-seq

- TopHat
  - Uses BowTie to map reads to genome
  - Analyzes mapping for evidence of splice junctions
  - Huge diversity of splice isoforms in higher eukaryotes



Front



